

# How to run the project :

On Linux :

***sudo apt install graphviz***

***cargo run***

I couldn't test on Windows, but graphviz has been ported for Windows, so it should work if installed.

main.rs contains an example of both way to define an `ExecutableGraph`

lib.rs contains the implementation of the thread pool

executable\_graph.rs contains the implementation of the `ExecutableGraph`

First a `ThreadPool` instance is necessary. We get one with the `ThreadPool::new` method, specifying how many threads we want.

Then there are two ways to generate an `ExecutableGraph`, as shown in *main.rs*

- using `ExecutableGraph::new` method, passing it an instance of `ThreadPool`
  - then calling `add_initial_node` and `add_node` methods to iteratively construct the graph
- using `ExecutableGraph::generate_random`, passing it
  - the number of nodes it will generate
  - the max number of parents a node can have
  - a list of arguments, which defines how many initial nodes are created and what to feed them with
  - a list of `Operation`, containing a closure and a name
  - an optional probability distribution for the operations
  - an instance of `ThreadPool`

Then using the `start` method to run the graph, and `show_graph` to print a representation of it. If too many nodes are used, it's probably wise not to use the `show_graph` method.

## Report

I am a beginner with Rust, but I had read the online book, and it was the occasion to practise.

### ***Comprehension of the problematic***

- at first, I thought of the parallelizability only in term of graph dependencies
  - every node that have the same parents are parallelizable
  - I thought I would create bundles of parallelizable nodes as clusters, and work with clusters instead of nodes and treat clusters sequentially
- then I realized the time factor matters A LOT
  - if one node inside a cluster took too much time, everything would wait for it
- in the end, I found a simpler way to tell which nodes to execute
  - those whose parents are done with their task
    - no need to test every nodes, just keeping track of which are possible candidates
  - once a candidate is done, track all its children instead
    - also update the edges with the result

- therefore, I needed a way to tell which nodes were finished
  - I used messages for the graph and the threads to communicate
    - each thread handles one node operation, and sends a message with the result and the identity of the node back to the graph

## Implementation

- since Rust doesn't support functions with arbitrary number of arguments, I used a vector
- I used to implement the Iterator trait to tell when there is no more node to execute, but traits are public, and it hadn't the behavior the user would expect
  - created `is_done` method instead
- tried to implement the graph itself with `Node` structure that hold references to other `Node` instances
  - big fight with the borrow checker
  - looked the implementation of petgraph to have an idea how they do it. They use indices.
  - ended up using petgraph to represent the graph
- `current_nodes` and `finished_nodes` used to be vectors
  - had to find elements inside it efficiently
  - ended up turning them into Hashmaps, as they can become really big, especially `finished_nodes`, who end up containing every nodes of the graph. Looking into the vector is  $O(n)$ , and  $O(1)$  with the hashmap
  - used to filter a lot the vectors to avoid adding nodes that were already there, used a hashset instead of hashmap
- Petgraph can generate .dot format files from the graph
  - but only if the types of the nodes and edges implement the `Display` trait, and the type `Option` doesn't
  - reconstruct an equivalent graph with String type for nodes and edges  
I guess I could have reimplemented my own `Option` type and implement the Display type on it instead

## Graph generation and execution

- the graph can only compute functions that have the same output type than the input  
 $f : X \times X \dots \times X \rightarrow X$  No verification is done on it. It can panic, for instance if the user defines a function that overflows an integer, or divide by zero.
- handle the acyclic constraint of the graph
  - first thought about trying to find circuits inside the graph everytime it's changed
  - then I realized that if instead of letting the user set the children of the nodes, we let him set the (already existing) parents, we can't end up with a cyclic graph
- how to tell which nodes are being executed in parallel ?
  - I know when I request the thread pool, but not exactly when the node is going to be executed
  - I could be notified when the thread starts its operation, via messaging
  - so I can know at an instant  $t$  which ones are being executed in parallel, by maintaining a list, but when do I print this list ? Every second ? I might be missing some if the operation is too fast

- in the end I simply print it everytime the main thread receive a message
- looked into openCL to try to execute on the GPU
  - but how do I pass an arbitrary rust function to the GPU ?
  - Apparently would require language translation  
how does TensorFlow does it ? Are all possible functions already predefined ?

### **Definition of initial nodes**

- initially thought I would look for every nodes which don't have parents
- since every node is an operation, they all have to have inputs, but inputs are defined by edges, then how do I define nodes that have no parent, like the entry points ?
- Thought they could be special type of nodes that would apply the identity function, but it's a waste to treat those as nodes to compute
  - ended up defining a special type of node called `initial\_node`, that would require arguments and that would be executed with those instead of the data defined in edges
  - other nodes require parents instead, they can't be orphan
    - therefore, the number of possible connected subgraphs is equal to the number of `initial nodes`